

MLP(X, Y, hidden, eta): trains a 2-layers MLP to fit X onto Y.

Args:

- X : network inputs
- Y : network targets
- η : learning rate
- max_epochs: maximum number of allowed epochs.
- $W_{init}^1, W_{init}^2, b_{init}^1, b_{init}^2$: initial weights and biases

Returns:

- W^1, W^2, b^1, b^2 : weights and biases of the connections between the input and the hidden layer and the hidden layer and the output layer.
- cost_hist: a history of cost values of different epochs in order of increasing epoch numbers

In order to calculate the weights we will use the backpropagation algorithm as discussed in the textbook. The MSE performance index is assumed and the algorithm trains for a 1000 epochs without checking for convergence. Note that the number of hidden layers is derived from initial weights and biases. We will also use momentum in the gradient descent step of the algorithm by adding a portion of the previous weights to the current update as follows:

$$V_t = \beta V_{t-1} + \eta \frac{\partial L}{\partial W}, V_0 = 0$$

$$W_{t+1} = W_t - V_t$$

Additionally, in order to obtain stable convergence, gradients updates are applied at the end of each epoch after averaging the gradients w.r.t. the entire dataset.

```
function [W1, W2, b1, b2, cost_hist] = MLP(X, Y, eta, max_epochs, W1_init, W2_init, b1_init, b2_init)
% cost history
cost_hist = zeros([max_epochs, 1]);
% momentum hyperparameter
beta = 0.9;
V_W1 = 0;
V_W2 = 0;
V_b1 = 0;
V_b2 = 0;
W1 = W1_init;
W2 = W2_init;
b1 = b1_init;
b2 = b2_init;
for i=1:max_epochs
    cost = 0;
    W1_grads = zeros(size(W1));
    W2_grads = zeros(size(W2));
    b1_grads = zeros(size(b1));
    b2_grads = zeros(size(b2));
    for k=1:length(X)
        p = X(:,k);
```

```

% target
t = Y(:,k);
% FORWARD PASS
% 1. Input Layer
n1 = W1*p + b1;
a1 = logsig(n1);

% 2. Hidden Layer
n2 = W2*a1 + b2;
a2 = logsig(n2);

% 3. Output and Error Calculation
loss = (t-a2)'*(t-a2); % squared error
cost = cost + loss;

% BACKWARD PASS
% 1. Calculating Sensitivities
% S2 = -2F_dot2(n2)*error
% S1 = F_dot1(n1)*W2'*S2
S2 = -2*diag(dlogsig(n2, a2), 0)*(t-a2);
S1 = diag(dlogsig(n1, a1), 0)*W2'*S2;

% 2. Weight Updates

W1_grads = W1_grads + S1*p';
W2_grads = W2_grads + S2*a1';

% 3. Bias Updates
b1_grads = b1_grads + S1;
b2_grads = b2_grads + S2;
end
V_W1 = beta*V_W1 + eta*W1_grads/length(X);
V_W2 = beta*V_W2 + eta*W2_grads/length(X);
V_b1 = beta*V_b1 + eta*b1_grads/length(X);
V_b2 = beta*V_b2 + eta*b2_grads/length(X);
W1 = W1 - V_W1;
W2 = W2 - V_W2;
b1 = b1 - V_b1;
b2 = b2 - V_b2;
if mod(i,100) == 0
    fprintf('Loss at the end of epoch %d: %f\n', i, cost/length(X));
end
cost_hist(i) = cost/length(X);
end
end

```